Model-call time and reported token use

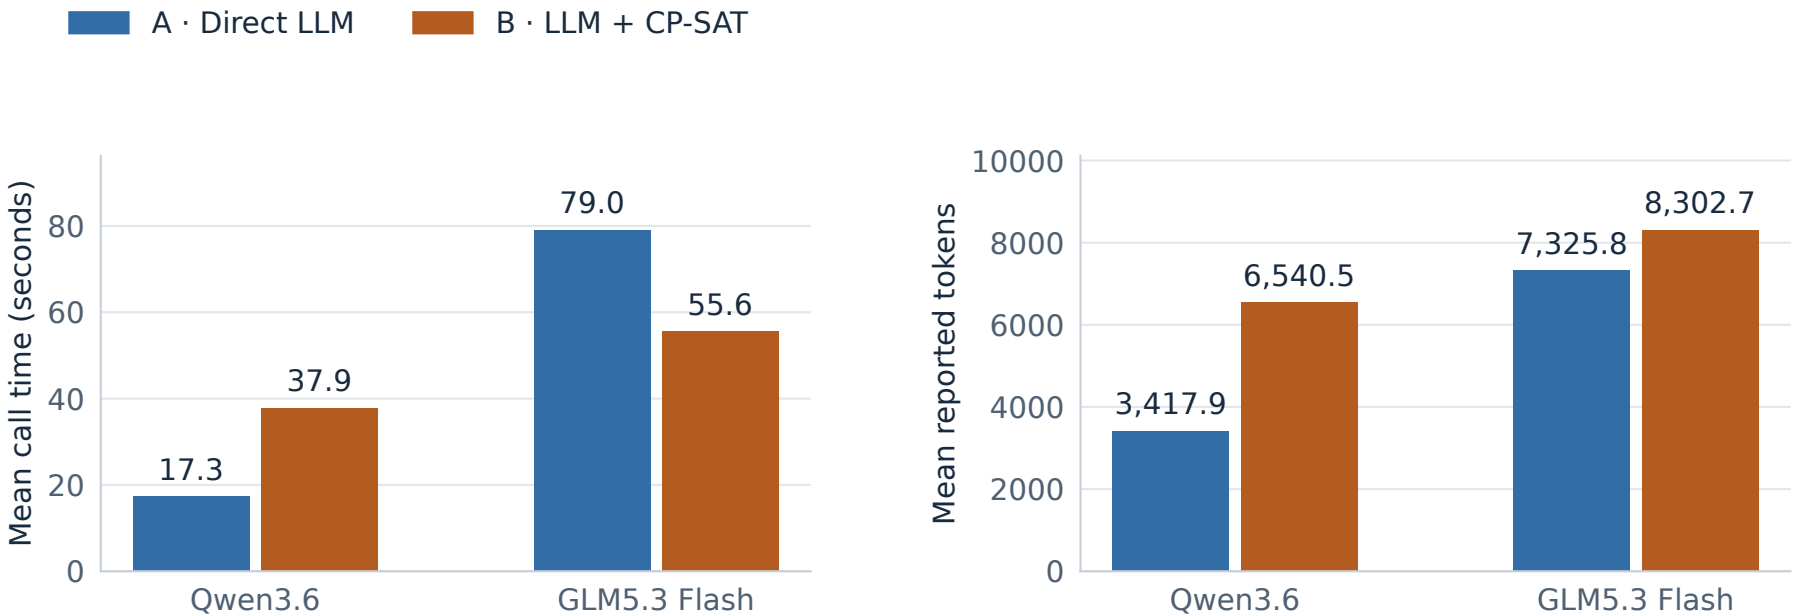

Time includes failed requests (100/100 per bar). Token coverage: Qwen A 96/100, B 98/100; GLM A 91/100, B 100/100. Missing usage is unknown. Shared-service timings do not isolate model speed; GLM A includes nine ~300-second HTTP 502s.